

R Basics

Why R?

- It's free!
- It runs on a variety of platforms including Windows, Unix and MacOS.
- It provides an unparalleled platform for programming new statistical methods in an easy and straightforward manner.
- It contains advanced statistical routines not yet available in other packages.
- It has state-of-the-art graphics capabilities.

How to download?

- Google it using R or CRAN
(Comprehensive R Archive Network)
- <http://www.r-project.org>
or
- <http://www.cran.r-project.org>

R Overview

R is a comprehensive statistical and graphical programming language and is a dialect of the S language:

1988 - S2: RA Becker, JM Chambers, A Wilks

1992 - S3: JM Chambers, TJ Hastie

1998 - S4: JM Chambers

R: initially written by Ross Ihaka and Robert Gentleman at Dep. of Statistics of U of Auckland, New Zealand during 1990s.

Since 1997: international “R-core” team of 15 people with access to common CVS archive.

R Overview

You can enter commands one at a time at the command prompt (>) or run a set of commands from a source file.

There is a wide variety of data types, including vectors (numerical, character, logical), matrices, data frames, and lists.

To quit R, use

```
>q()
```

R Overview

Most functionality is provided through built-in and user-created functions and all data objects are kept in memory during an interactive session.

Basic functions are available by default. Other functions are contained in packages that can be attached to a current session as needed

R Interface

Start the R system, the main window (RGui)
with a sub window (R Console) will appear
In the 'Console' window the cursor is waiting
for you to type in some R commands.

Your First R Session

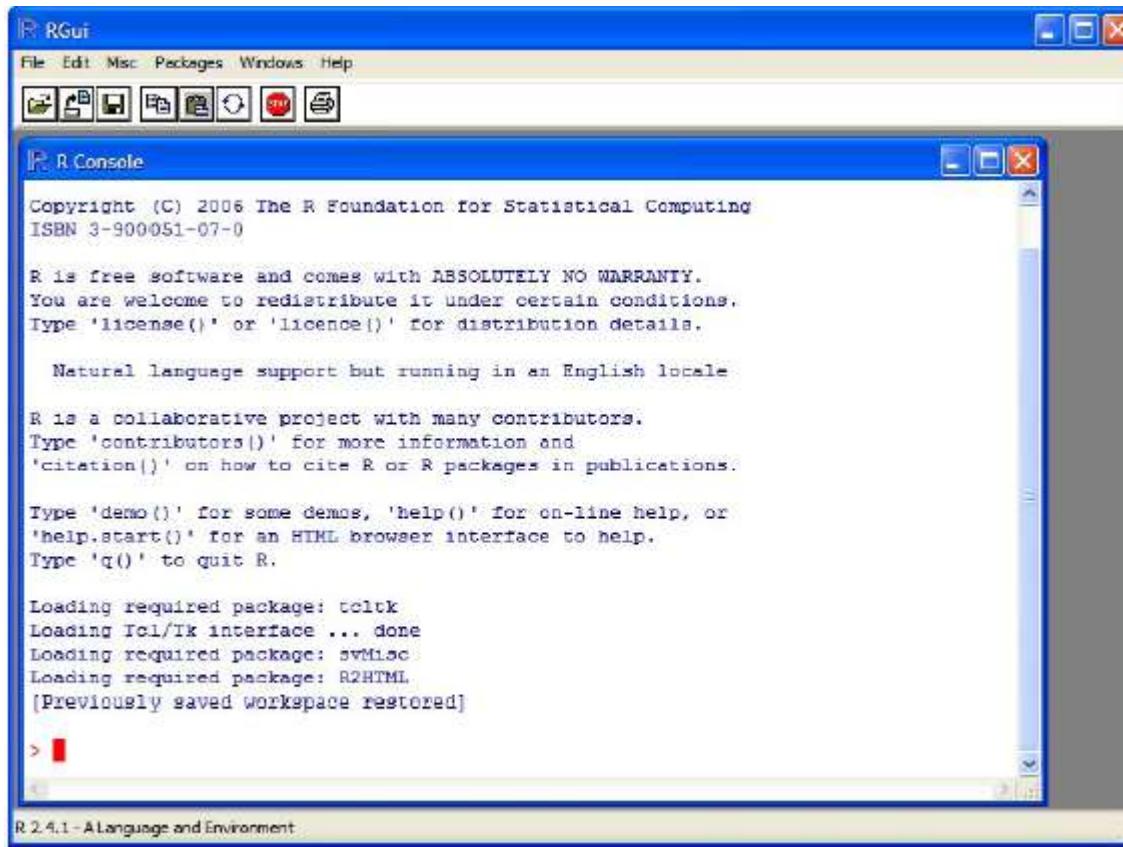
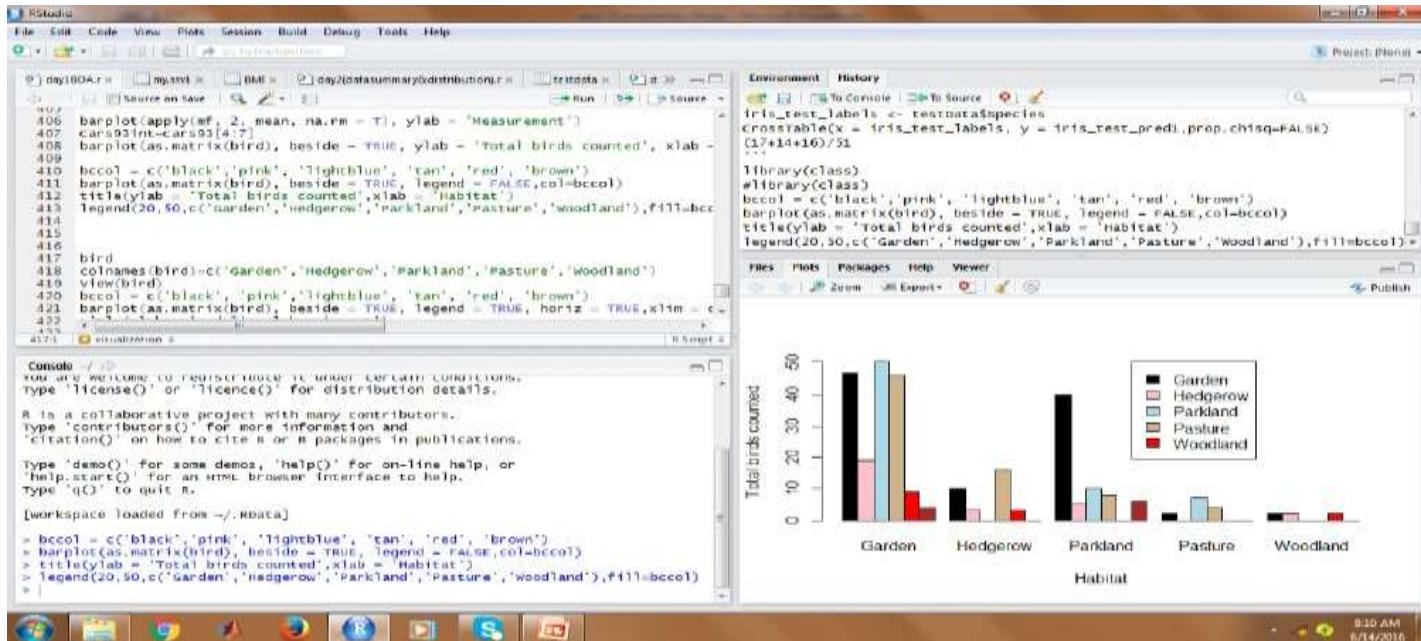


Figure 1.1: The R system on Windows

- **R Studio:** R Studio is an Integrated Development Environment (IDE) for R Language with advanced and more user-friendly GUI. R Studio allows the user to run R in a more user-friendly environment. It is open-source (i.e.free) and available at <http://www.rstudio.com/>.



R Introduction

- Results of calculations can be stored in objects using the assignment operators:
 - An arrow (<-) formed by a smaller than character and a hyphen without a space!
 - The equal character (=).

R Introduction

- These objects can then be used in other calculations. To print the object just enter the name of the object. There are some restrictions when giving an object a name:
 - Object names cannot contain 'strange' symbols like !, +, -, #.
 - A dot (.) and an underscore () are allowed, also a name starting with a dot.
 - Object names can contain a number but cannot start with a number.
 - R is case sensitive, X and x are two different objects, as well as temp and tempP.

An example

```
> # An example
> x <- c(1:10)
> x[(x>8) | (x<5)]
> # yields 1 2 3 4 9 10
> # How it works
> x <- c(1:10)
> x
>1 2 3 4 5 6 7 8 9 10
> x > 8
> F F F F F F F T T
> x < 5
> T T T T F F F F F F
> x > 8 | x < 5
> T T T T F F F F T T
> x[c(T,T,T,T,F,F,F,F,T,T)]
> 1 2 3 4 9 10
```

R Introduction

- To list the objects that you have in your current R session use the function ls or the function objects.

```
> ls()  
[1] "x" "y"
```
- So to run the function ls we need to enter the name followed by an opening (and a closing). Entering only ls will just print the object, you will see the underlying R code of the the function ls. Most functions in R accept certain arguments. For example, one of the arguments of the function ls is pattern. To list all objects starting with the letter x:

```
> x2 = 9  
> y2 = 10  
> ls(pattern="x")  
[1] "x2"  
>x1 <- 1 # Create data objects  
>x2 <- 2  
>x3 <- 3  
>ls()  
[1] "x1" "x2" "x3"
```

R Warning !

R is a case sensitive language.
FOO, Foo, and foo are three different objects

R Introduction

```
> x = sin(9)/75
> y = log(x) + x^2
> x
[1] 0.005494913
> y
[1] -5.203902
> m <- matrix(c(1,2,4,1), ncol=2)
> m
> [,1] [,2]
[1,] 1 4
[2,] 2 1
> solve(m)
      [,1]      [,2]
[1,] -0.1428571 0.5714286
[2,] 0.2857143 -0.1428571
```

R Workspace

```
getwd() # print the current working directory
```

```
ls() # list the objects in the current workspace
```

```
setwd(mydirectory) # change to mydirectory
```

```
setwd("c:/docs/mydir")
```

R Help

Once R is installed, there is a comprehensive built-in help system. At the program's command prompt you can use any of the following:

```
help.start() # general help  
help(foo)   # help about function foo  
?foo        # same thing  
apropos("foo") # list all functions containing string foo  
example(foo) # show an example of function foo
```

R Datasets

R comes with a number of sample datasets that you can experiment with. Type

> data()

to see the available datasets. The results will depend on which packages you have loaded.

Type

help(datasetname)

for details on a sample dataset.

R Packages

- One of the strengths of R is that the system can easily be extended. The system allows you to write new functions and package those functions in a so called 'R package' (or 'R library'). The R package may also contain other R objects, for example data sets or documentation. There is a lively R user community and many R packages have been written and made available on CRAN for other users. Just a few examples, there are packages for portfolio optimization, drawing maps, exporting objects to html, time series analysis, spatial statistics and the list goes on and on.

R Packages

- When you download R, already a number (around 30) of packages are downloaded as well. To use a function in an R package, that package has to be attached to the system. When you start R not all of the downloaded packages are attached, only seven packages are attached to the system by default. You can use the function search to see a list of packages that are currently attached to the system, this list is also called the search path.

```
> search()  
[1] ".GlobalEnv" "package:stats" "package:graphics"  
[4] "package:grDevices" "package:datasets" "package:utils"  
[7] "package:methods" "Autoloads" "package:base"
```

R Packages

To attach another package to the system you can use the menu or the library function. Via the menu:

Select the 'Packages' menu and select 'Load package...', a list of available packages on your system will be displayed. Select one and click 'OK', the package is now attached to your current R session. Via the library function:

```
> library(MASS)
> shoes
$A
[1] 13.2 8.2 10.9 14.3 10.7 6.6 9.5 10.8 8.8 13.3
$B
[1] 14.0 8.8 11.2 14.2 11.8 6.4 9.8 11.3 9.3 13.6
```

R Packages

- The function `library` can also be used to list all the available libraries on your system with a short description. Run the function without any arguments

```
> library()
```

Packages in library 'C:/PROGRA~1/R/R-25~1.0/library':

base	The R Base Package
Boot	Bootstrap R (S-Plus) Functions (Canty)
class	Functions for Classification
cluster	Cluster Analysis Extended Rousseeuw et al.
codetools	Code Analysis Tools for R
datasets	The R Datasets Package
DBI	R Database Interface
foreign	Read Data Stored by Minitab, S, SAS, SPSS, Stata, Systat, dBase, ...
graphics	The R Graphics Package

Lecture 2: Data Input

Data Types

R has a wide variety of data types including scalars, vectors (numerical, character, logical), matrices, dataframes, and lists.

Vectors

- A vector in R is a one-dimensional array that can contain elements of the **same data type**.
- **To create a vector, you can use the c() function.**

For example:

- Numeric vector:
1) my_vector <- c(1, 2, 3, 4, 5)
2)a <- c(1,2,5.3,6,-2,4)
- Character vector:
1) fruits <- c("apple", "banana", "cherry")
2) b <- c("one","two","three")
- logical vector:c <- c(TRUE,TRUE,TRUE,FALSE,TRUE,FALSE)
- Refer to elements of a vector using subscripts.
a[c(2,4)] # 2nd and 4th elements of vector

- **any():** The **any()** function returns TRUE if at least one element in a logical vector is TRUE.

For example:

```
logical_vector <- c(TRUE, FALSE, FALSE)  
result <- any(logical_vector) # Returns TRUE
```

- **all():** The **all()** function returns TRUE if all elements in a logical vector are TRUE.

For example

```
logical_vector <- c(TRUE, TRUE, TRUE)  
result <- all(logical_vector) # Returns TRUE
```

- **which():** The **which()** function returns the indices of elements in a logical vector that are TRUE.

For example:

```
logical_vector <- c(TRUE, FALSE, TRUE, TRUE, FALSE)  
indices <- which(logical_vector) # Returns indices 1, 3, 4
```

- To sort a vector in descending order in R, you can use the **sort() function with the decreasing parameter set to TRUE**.

For example:

1. `vec <- c(5, 2, 8, 1, 3)`
2. `sorted_vec <- sort(vec, decreasing = TRUE)` will give you a vector sorted in descending order.

- **length: The length() function is used to determine the length of a vector, which is the number of elements it contains.**

For example:

```
my_vector <- c(4, 7, 2, 9, 5)
```

```
vector_length <- length(my_vector) # Returns the length of the vector, which is 5
```

- **seq:** The seq() function in R is used to create sequences of numbers. It can be used to generate a sequence from a starting value to an ending value, with a specified increment.

For example:

```
sequence <- seq(1, 10, by = 2) # Creates a sequence: 1, 3, 5, 7, 9
```

- **rep:** The rep() function is used to replicate elements in a vector.

It takes two main arguments: the vector to be replicated and the number of times to repeat it.

For example:

```
vector <- c(1, 2, 3)
```

```
repeated_vector <- rep(vector, times = 3)
```

```
# Repeats the vector three times: 1, 2, 3, 1, 2, 3, 1, 2, 3
```

Matrices

All columns in a matrix must have the same mode(numeric, character, etc.) and the same length.

To create a matrix:

The general format is

```
mymatrix <- matrix(vector, nrow=r, ncol=c,  
  byrow=FALSE,dimnames=list(char_vector_rownames,  
  char_vector_colnames))
```

byrow=TRUE indicates that the matrix should be filled by rows.

byrow=FALSE indicates that the matrix should be filled by columns
(the default).

dimnames provides optional labels for the columns and rows.

- cbind and rbind are functions in R used to combine vectors or matrices into a matrix.
- cbind combines objects **by columns**,
- rbind combines objects **by rows**.
- **For example:**

```
x <- c(1, 2, 3)
```

```
y <- c(4, 5, 6)
```

combined <- cbind(x, y) will create a 3x2 matrix with x and y as columns.

combined <- rbind(x, y) will create a 3x2 matrix with x and y as rows.

- The `diag()` function in R is used to create a diagonal matrix or extract the diagonal elements from a matrix.

For example:

- To create a diagonal matrix: `diag_matrix <- diag(1:4)`
- To extract the diagonal elements from a matrix:

```
mat<-matrix(1:9,nrow=3)
```

```
diagonal_elements <- diag(mat)
```

- To find the transpose and inverse of a matrix in R:

```
mat<-matrix(1:9,nrow=3)
```

- Transpose: `t(mat)` #where mat is your matrix.
- Inverse: `solve(mat)` #where mat is a square, invertible matrix.

- To find the dimension of a matrix in R, you can use the **dim()** function.
- **For example:**

```
mat <- matrix(1:6, nrow = 2, ncol = 3)
```

```
dim(mat) #will return the dimensions as 2 3,  
indicating it's a 2x3 matrix.
```

Example

```
mat <- matrix(c(4.3, 3.1, 8.2, 8.2, 3.2, 0.9, 1.6, 6.5), nrow = 4, ncol = 2, byrow = TRUE)
```

Output

```
[,1] [,2]  
[1,] 4.3 3.1  
[2,] 8.2 8.2  
[3,] 3.2 0.9  
[4,] 1.6 6.5
```

- To omit and overwrite elements in a matrix, you can simply assign new values to the elements you want to modify. For example, to omit the element in the second row and third column and overwrite an element:

```
mat <- matrix(1:9, nrow = 3)
```

```
mat[2, 3] <- NA # Omit the element
```

```
mat[1, 1] <- 99 # Overwrite an element
```

Matrix Operations and Algebra

1) Matrix Transpose

For any $m \times n$ matrix A , its transpose, A^\top , is then $m \times n$ matrix obtained by writing either its columns as rows or its rows as columns

$$\text{If } A = \begin{bmatrix} 2 & 5 & 2 \\ 6 & 1 & 4 \end{bmatrix}, \text{ then } A^\top = \begin{bmatrix} 2 & 6 \\ 5 & 1 \\ 2 & 4 \end{bmatrix}.$$

Matrix Operations and Algebra

2) IdentityMatrix

The identity matrix written as I_m is a particular kind of matrix used in mathematics. It's a square $m \times m$ matrix with ones on the diagonal and zeros elsewhere.

$$I_3 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

```
R> A <- diag(x=3)
R> A
      [,1] [,2] [,3]
[1,]    1    0    0
[2,]    0    1    0
[3,]    0    0    1
```

3) Scalar Multiple of a Matrix

A scalar value is just a single, univariate value. Multiplication of any matrix A by a scalar value a results in a matrix in which every individual element is multiplied by a.

$$2 \times \begin{bmatrix} 2 & 5 & 2 \\ 6 & 1 & 4 \end{bmatrix} = \begin{bmatrix} 4 & 10 & 4 \\ 12 & 2 & 8 \end{bmatrix}$$

```
R> A <- rbind(c(2,5,2),c(6,1,4))
R> a <- 2
R> a*A
     [,1] [,2] [,3]
[1,]    4   10    4
[2,]   12     2    8
```

4) Matrix addition and subtraction

Addition or subtraction of two matrices of equal size is also performed in an element-wise fashion. Corresponding elements are added or subtracted from one another, depending on the operation.

$$\begin{bmatrix} 2 & 6 \\ 5 & 1 \\ 2 & 4 \end{bmatrix} - \begin{bmatrix} -2 & 8.1 \\ 3 & 8.2 \\ 6 & -9.8 \end{bmatrix} = \begin{bmatrix} 4 & -2.1 \\ 2 & -7.2 \\ -4 & 13.8 \end{bmatrix}$$

```
R> A <- cbind(c(2,5,2),c(6,1,4))
R> A
      [,1] [,2]
[1,]    2    6
[2,]    5    1
[3,]    2    4
R> B <- cbind(c(-2,3,6),c(8.1,8.2,-9.8))
R> B
      [,1] [,2]
[1,]   -2   8.1
[2,]    3   8.2
[3,]    6  -9.8
R> A-B
```

Matrix

```
      [,1] [,2]
[1,]    4   -2.1
[2,]    2   -7.2
[3,]   -4   13.8
```

5) Matrix Multiplication

In order to multiply two matrices A and B of size m X n and p X q, it must be true that n=p. The resulting matrix A . B will have the size m X q. The elements of the product are computed in a row-by-column fashion, where the value at position(AB)_{i,j} is computed by element-wise multiplication of the entries in row i of A by the entries in column j of B, summing the result.

$$\begin{aligned} & \begin{bmatrix} 2 & 5 & 2 \\ 6 & 1 & 4 \end{bmatrix} \cdot \begin{bmatrix} 3 & -3 \\ -1 & 1 \\ 1 & 5 \end{bmatrix} \\ &= \begin{bmatrix} 2 \times 3 + 5 \times (-1) + 2 \times 1 & 2 \times (-3) + 5 \times (1) + 2 \times 5 \\ 6 \times 3 + 1 \times (-1) + 4 \times 1 & 6 \times (-3) + 1 \times (1) + 4 \times 5 \end{bmatrix} \\ &= \begin{bmatrix} 3 & 9 \\ 21 & 3 \end{bmatrix} \end{aligned}$$

6) Matrix Inversion

Some square matrices can be inverted. The inverse of a matrix A is denoted A^{-1} . An invertible matrix satisfies the following equation:

$$A A^{-1} = I_m$$

$$\begin{bmatrix} 3 & 1 \\ 4 & 2 \end{bmatrix}^{-1} = \begin{bmatrix} 1 & -0.5 \\ -2 & 1.5 \end{bmatrix}$$

```
R> A <- matrix(data=c(3,4,1,2), nrow=2, ncol=2)
R> A
     [,1] [,2]
[1,]    3    1
[2,]    4    2
R> solve(A)
     [,1] [,2]
[1,]    1 -0.5
[2,]   -2   1.5
```

Arrays

R Arrays consist of all elements of the same data type. Vectors are supplied as input to the function and then create an array based on the number of dimensions.

You can create arrays in R using the `array()` function. An array is a multi-dimensional structure that can hold data of the same type.

- **Syntax:**

```
array(data, dim = (nrow, ncol, nmat), dimnames=names)
```

where

- `nrow`: Number of rows
- `ncol` : Number of columns
- `nmat`: Number of matrices of dimensions `nrow * ncol`
- `dimnames` : Default value = `NULL`.
- Here's an example of creating a $3 \times 3 \times 3$ array:

```
data <- 1:27 # Vector with 27 elements
```

```
my_array <- array(data, dim = c(3, 3, 3))
```

Multidimensional Arrays

- Just as a matrix (a “rectangle” of elements) is the result of increasing the dimension of a vector (a “line” of elements), the dimension of a matrix can be increased to get more complex data structures.
- three-dimensional array can be considered to be a collection of equally dimensioned matrices, providing you with a rectangular prism of elements.
- Figure 3-3 illustrates a three-row, four-column, two-layer (3 4 2) array.

Multidimensional Arrays

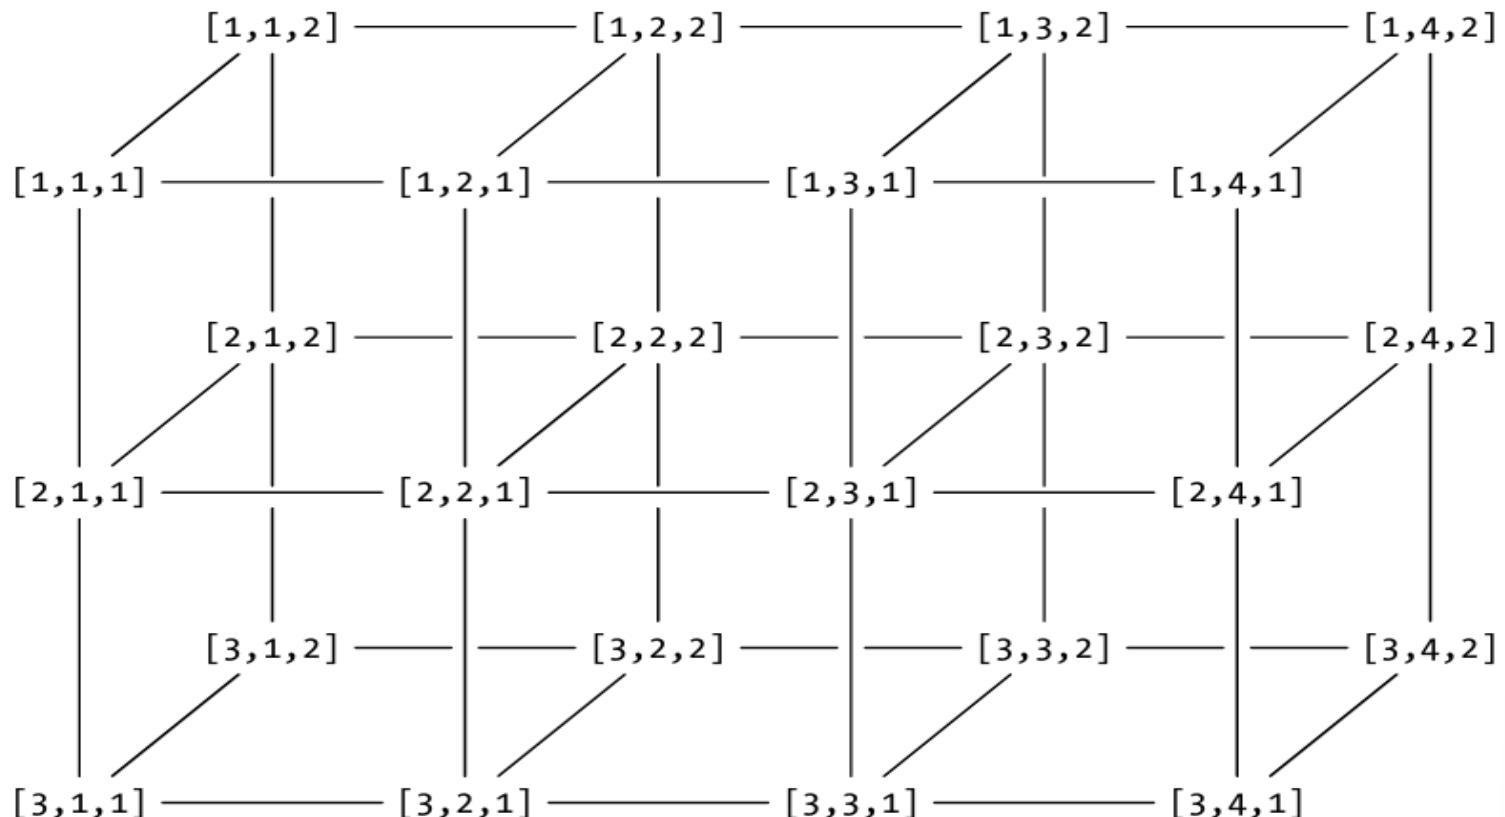


Figure 3-3: A conceptual diagram of a $3 \times 4 \times 2$ array. The index of each element is given at the corresponding position. These indexes are provided in the strict order of [row, column, layer].

Multidimensional Arrays

Definition

To create these data structures in R, use the array function and specify the individual elements in the data argument as a vector. Then specify size in the dim argument as another vector with a length corresponding to the number of dimensions.

Note that array fills the entries of each layer with the elements in data in a strict column-wise fashion, starting with the first layer.

Multidimensional Arrays

```
R> AR <- array(data=1:24, dim=c(3,4,2))
R> AR
, , 1
```

	[,1]	[,2]	[,3]	[,4]
[1,]	1	4	7	10
[2,]	2	5	8	11
[3,]	3	6	9	12

```
, , 2
```

	[,1]	[,2]	[,3]	[,4]
[1,]	13	16	19	22
[2,]	14	17	20	23
[3,]	15	18	21	24

Subsets, Extractions, and Replacements

Suppose you want the second row of the second layer of the previously created array AR. You just enter these exact dimensional locations of AR in square brackets.

```
R> AR[2,,2]
```

```
[1] 14 17 20 23
```

The desired elements have been extracted as a vector of length 4. If you want specific elements from this vector, say the third and first, in that order, you can call the following:

```
R> AR[2,c(3,1),2]
```

```
[1] 20 14
```

Subsets, Extractions, and Replacements

- To extract the first rows of both layers of AR, you enter this:

```
R> AR[1,,]
```

```
[,1] [,2]  
[1,] 1 13  
[2,] 4 16  
[3,] 7 19  
[4,] 10 22
```

Data frames

- In R, a data frame is a two-dimensional structure that stores data in rows and columns, similar to a table in a database or a spreadsheet. Each column in a data frame can have a different data type, and it can store various types of data such as numeric, character, logical, etc.
- # Creating the data frame

```
df <- data.frame(  
  Person = c("Peter", "Lois", "Meg", "Chris", "Stewie"),  
  Age = c(42, 40, 17, 14, 1),  
  Sex = c("M", "F", "F", "M", "M"))  
  
# Viewing the created data frame  
print(df)
```

OUTPUT

```
Person Age Sex  
1 Peter 42 M  
2 Lois 40 F  
3 Meg 17 F  
4 Chris 14 M  
5 Stewie 1 M
```

There are a variety of ways to identify the elements of a dataframe

```
myframe[3:5] # columns 3,4,5 of dataframe
```

```
myframe[c("ID","Age")] # columns ID and Age from dataframe
```

```
myframe$X1 # variable x1 in the dataframe
```

Data frames

Lists

An ordered collection of objects (components). A list allows you to gather a variety of (possibly unrelated) objects under one name.

```
# example of a list with 4 components -
```

```
# a string, a numeric vector, a matrix, and a scalar
```

```
w <- list(name="Fred", mynumbers=a, mymatrix=y, age=5.3)
```

```
# example of a list containing two lists
```

```
v <- c(list1,list2)
```

Lists

list in R is a data structure that can hold a collection of objects of different types.

For

- example:

```
my_list <- list(name = "John", age = 30, scores = c(85, 90, 78))
```

- Identify elements of a list using the [[[]]] convention.

mylist[[2]] # 2nd component of the list

- List slicing can be done using indexing and subsetting.

- For example:-

```
my_list <- list("apple", "orange", "banana", "grape", "melon")
```

```
subset <- my_list[2:4] # Extracting a subset of elements from  
# index 2 to 4 (inclusive)
```

```
print(subset)
```

Factors

A factor is a data structure in R used to represent categorical variables. Factors are used to store data that can take on a limited, fixed number of values, known as levels. Tell **R** that a variable is **nominal** by making it a factor.

The factor stores the nominal values as a vector of integers in the range [1... k] (where k is the number of unique values in the nominal variable), and an internal vector of character strings (the original values) mapped to these integers.

```
# variable gender with 20 "male" entries and # 30 "female" entries  
gender <- c(rep("male",20), rep("female", 30))  
gender <- factor(gender) # stores gender as 20 1s and 30 2s and associates#  
# 1=female, 2=male internally (alphabetically)  
# R now treats gender as a nominal variable  
summary(gender)
```

OUTPUT

female male

30 20

- The levels() function in R is used with factors to get the levels distinct categories of a factor variable.
- For example:

```
gender <- factor(c("Male", "Female", "Male",
"Male", "Female"))
```

```
gender_levels <- levels(gender) #will return a
vector of levels: "Female" "Male".
```

- **cat:** The `cat()` function in R is used to concatenate and print values together. It prints the values without any separators by default.

```
cat("Hello", "World")
```

Output: Hello World

- **paste:** The `paste()` function is used to concatenate strings or other objects with a specified separator.

```
result <- paste("Hello", "World", sep = ", ")
```

Result: "Hello, World"

Useful Functions

```
length(object) # number of elements or components  
str(object)  # structure of an object  
class(object) # class or type of an object  
names(object) # names  
c(object,object,...) # combine objects into a vector  
cbind(object, object, ...) # combine objects as columns  
rbind(object, object, ...) # combine objects as rows  
ls()      # list current objects  
rm(object) # delete an object  
newobject <- edit(object) # edit copy and save a newobject  
fix(object)       # edit in place
```

Viewing Data

There are a number of functions for listing the contents of an object or dataset.

```
# list objects in the working environment  
ls()  
  
# list the variables in mydata  
names(mydata)  
  
# list the structure of mydata  
str(mydata)  
  
# list levels of factor v1 in mydata  
levels(mydata$v1)  
  
# dimensions of an object  
dim(object)
```

Viewing Data

There are a number of functions for listing the contents of an object or dataset.

```
# class of an object (numeric, matrix, dataframe, etc)
  class(object)

# print mydata
  mydata

# print first 10 rows of mydata
  head(mydata, n=10)

# print last 5 rows of mydata
  tail(mydata, n=5)
```

Missing Data

- In R, missing values are represented by the symbol **NA** (not available) .
- Impossible values (e.g., dividing by zero) are represented by the symbol **NaN** (not a number).
- Unlike SAS(Statistical Analysis System), R uses the same symbol for character and numeric data.

Testing for Missing Values

```
is.na(x) # returns TRUE if x is missing
```

```
y <- c(1,2,3,NA)
```

```
is.na(y) # returns a vector (F F F T)
```

Missing Data

Recoding Values to Missing

```
# recode 99 to missing for variable v1  
# select rows where v1 is 99 and recode column v1  
mydata[mydata$v1==99,"v1"] <- NA
```

Excluding Missing Values from Analyses

Arithmetic functions on missing values yield missing values.

```
x <- c(1,2,NA,3)  
mean(x)      # returns NA  
mean(x, na.rm=TRUE) # returns 2
```

Missing Data

The function **complete.cases()** returns a logical vector indicating which cases are complete.

```
# list rows of data that have missing values  
mydata[!complete.cases(mydata),]
```

The function **na.omit()** returns the object with listwise deletion of missing values.

```
# create new dataset without missing data  
newdata <- na.omit(mydata)
```

Date Values

**Dates are represented as the number of days since 1970-01-01,
with negative values for earlier dates.**

```
# use as.Date( ) to convert strings to dates  
mydates <- as.Date(c("2007-06-22", "2004-02-13"))  
# number of days between 6/22/07 and 2/13/04  
days <- mydates[1] - mydates[2]
```

Sys.Date() returns today's date.

Date() returns the current date and time.

Date Values

The following symbols can be used with the `format()` function to print dates.

Symbol	Meaning	Example
<code>%d</code>	day as a number (0-31)	01-31
<code>%a</code>	abbreviated weekday	Mon
<code>%A</code>	unabbreviated weekday	Monday
<code>%m</code>	month (00-12)	00-12
<code>%b</code>	abbreviated month	Jan
<code>%B</code>	unabbreviated month	January
<code>%y</code>	2-digit year	07
<code>%Y</code>	4-digit year	2007

Date Values

```
# print today's date
today <- Sys.Date()
format(today, format="%B %d %Y")
"June 20 2007"
```

Lecture 3: Data Manipulation

Introduction

Each of these activities usually involve the use of R's built-in [operators](#) (arithmetic and logical) and [functions](#) (numeric, character, and statistical). Additionally, you may need to use [control structures](#) (if-then, for, while, switch) in your programs and/or create your [own functions](#). Finally you may need to [convert](#) variables or datasets from one type to another (e.g. numeric to character or matrix to dataframe).

Arithmetic Operators

Operator	Description
+	addition
-	subtraction
*	multiplication
/	division
^ or **	exponentiation
x %% y	modulus (x mod y) 5%%2 is 1
x %/% y	integer division 5%/%2 is 2

Logical Operators

Operator	Description
<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
==	exactly equal to
!=	not equal to
!x	Not x
x y	x OR y
x & y	x AND y
isTRUE(x)	test if x is TRUE

Control Structures

- R has the standard control structures you would expect. **expr** can be multiple (compound) statements by enclosing them in braces { }. It is more efficient to use built-in functions rather than control structures whenever possible.

Control Structures

if-else

```
if (condition) {  
    do_this  
} else {  
    do_that  
}
```

Example

```
x <- 10
```

```
if (x > 5) {  
    print("x is greater than 5")  
} else {  
    print("x is 5 or less")  
}
```

Control Structures

for

```
for (variable in sequence) {  
    do_this  
}
```

Example

```
for (i in 1:3) {  
    print(i)  
}
```

Control Structures

while

```
while (condition) {  
    do_this  
}
```

Example

```
x <- 1  
while (x <= 3) {  
    print(x)  
    x <- x + 1  
}
```

Control Structures

switch

```
switch(expression,  
       case1 = result1,  
       case2 = result2,  
       ...)
```

Example

```
x <- 2  
switch(x,  
       "one",  
       "two",  
       "three")
```

Control Structures

ifelse

ifelse(*test*,*yes*,*no*)

Example

x <- c(4, 6, 8)

```
result <- ifelse(x > 5, "big", "small")
print(result)
```

Control Structures

- # transpose of a matrix
a poor alternative to built-in t() function

```
mytrans <- function(x) {
  if (!is.matrix(x)) {
    warning("argument is not a matrix: returning NA")
    return(NA_real_)
  }
  y <- matrix(1, nrow=ncol(x), ncol=nrow(x))
  for (i in 1:nrow(x)) {
    for (j in 1:ncol(x)) {
      y[j,i] <- x[i,j]
    }
  }
  return(y)
}
```

Control Structures

- # try it
z <- matrix(1:10, nrow=5, ncol=2)
tz <- mytrans(z)

R built-in functions

Almost everything in R is done through functions. Here I'm only referring to numeric and character functions that are commonly used in creating or recoding variables.

Note that while the examples on this page apply functions to individual variables, many can be applied to vectors and matrices as well.

Numeric Functions

Function	Description
abs(x)	absolute value
sqrt(x)	square root
ceiling(x)	ceiling(3.475) is 4
floor(x)	floor(3.475) is 3
trunc(x)	trunc(5.99) is 5
round(x, digits=n)	round(3.475, digits=2) is 3.48
signif(x, digits=n)	signif(3.475, digits=2) is 3.5
cos(x), sin(x), tan(x)	also $\text{acos}(x)$, $\text{cosh}(x)$, $\text{acosh}(x)$, etc.
log(x)	natural logarithm
log10(x)	common logarithm
exp(x)	e^x

Character Functions

Function	Description
substr(x, start=n1, stop=n2)	Extract or replace substrings in a character vector. <code>x <- "abcdef"</code> <code>substr(x, 2, 4) is "bcd"</code> <code>substr(x, 2, 4) <- "22222" is "a222ef"</code>
grep(pattern, x , ignore.case=FALSE, fixed=FALSE)	Search for <i>pattern</i> in <i>x</i> . If <i>fixed</i> =FALSE then <i>pattern</i> is a regular expression . If <i>fixed</i> =TRUE then <i>pattern</i> is a text string. Returns matching indices. <code>grep("A", c("b","A","c"), fixed=TRUE) returns 2</code>
sub(pattern, replacement, x, ignore.case =FALSE, fixed=FALSE)	Find <i>pattern</i> in <i>x</i> and replace with <i>replacement</i> text. If <i>fixed</i> =FALSE then <i>pattern</i> is a regular expression. If <i>fixed</i> = T then <i>pattern</i> is a text string. <code>sub("\\s",".", "Hello There") returns "Hello.There"</code>
strsplit(x, split)	Split the elements of character vector <i>x</i> at <i>split</i> . <code>strsplit("abc", "")</code> returns 3 element vector "a","b","c"
paste(..., sep="")	Concatenate strings after using <i>sep</i> string to seperate them. <code>paste("x",1:3,sep="")</code> returns <code>c("x1","x2" "x3")</code> <code>paste("x",1:3,sep="M")</code> returns <code>c("xM1","xM2" "xM3")</code> <code>paste("Today is", date())</code>
toupper(x)	Uppercase
tolower(x)	Lowercase

Stat/Prob Functions

- The following table describes functions related to probability distributions. For random number generators below, you can use `set.seed(1234)` or some other integer to create reproducible pseudo-random numbers.

Other Useful Functions

Function	Description
seq (<i>from</i> , <i>to</i> , <i>by</i>)	generate a sequence indices <- seq(1,10,2) #indices is c(1, 3, 5, 7, 9)
rep (<i>x</i> , <i>ntimes</i>)	repeat <i>x</i> <i>n</i> times y <- rep(1:3, 2) # y is c(1, 2, 3, 1, 2, 3)
cut (<i>x</i> , <i>n</i>)	divide continuous variable in factor with <i>n</i> levels y <- cut(x, 5)

Function	Description
dnorm(<i>x</i>)	normal density function (by default m=0 sd=1) # plot standard normal curve <pre>x <- pretty(c(-3,3), 30) y <- dnorm(x) plot(x, y, type='l', xlab="Normal Deviate", ylab="Density", yaxs="i")</pre>
pnorm(<i>q</i>)	cumulative normal probability for <i>q</i> (area under the normal curve to the right of <i>q</i>) <i>pnorm(1.96)</i> is 0.975
qnorm(<i>p</i>)	normal quantile. value at the <i>p</i> percentile of normal distribution <i>qnorm(.9)</i> is 1.28 # 90th percentile
rnorm(<i>n</i>, m=0, sd=1)	<i>n</i> random normal deviates with mean <i>m</i> and standard deviation <i>sd</i> . <pre>#50 random normal variates with mean=50, sd=10 x <- rnorm(50, m=50, sd=10)</pre>
dbinom(<i>x</i>, size, prob) pbinom(<i>q</i>, size, prob) qbinom(<i>p</i>, size, prob) rbinom(<i>n</i>, size, prob)	binomial distribution where size is the sample size and prob is the probability of a heads (pi) <pre># prob of 0 to 5 heads of fair coin out of 10 flips dbinom(0:5, 10, .5)</pre> <pre># prob of 5 or less heads of fair coin out of 10 flips pbinom(5, 10, .5)</pre>
dpois(<i>x</i>, lamda) ppois(<i>q</i>, lamda) qpois(<i>p</i>, lamda) rpois(<i>n</i>, lamda)	poisson distribution with m=std=lamda <pre>#probability of 0,1, or 2 events with lamda=4 dpois(0:2, 4)</pre> <pre># probability of at least 3 events with lamda=4 1- ppois(2,4)</pre>
dunif(<i>x</i>, min=0, max=1) punif(<i>q</i>, min=0, max=1) qunif(<i>p</i>, min=0, max=1) runif(<i>n</i>, min=0, max=1)	uniform distribution, follows the same pattern as the normal distribution above. <pre>#10 uniform random variates x <- runif(10)</pre>

Function	Description
mean(<i>x</i>, trim=0, na.rm=FALSE)	mean of object x # trimmed mean, removing any missing values and # 5 percent of highest and lowest scores mx <- mean(x,trim=.05,na.rm=TRUE)
sd(<i>x</i>)	standard deviation of object(x). also look at var(x) for variance and mad(x) for median absolute deviation.
median(<i>x</i>)	median
quantile(<i>x, probs</i>)	quantiles where x is the numeric vector whose quantiles are desired and probs is a numeric vector with probabilities in [0,1]. # 30th and 84th percentiles of x y <- quantile(x, c(.3,.84))
range(<i>x</i>)	range
sum(<i>x</i>)	sum
diff(<i>x, lag=1</i>)	lagged differences, with lag indicating which lag to use
min(<i>x</i>)	minimum
max(<i>x</i>)	maximum
scale(<i>x, center=TRUE, scale=TRUE</i>)	column center or standardize a matrix.

Lecture 4: Non-Numeric Values

Logical Values

Logical values (also simply called **logicals**) are based on a simple premise: a logical-valued object can only be either TRUE or FALSE. These can be interpreted as yes/no, one/zero, satisfied/not satisfied, and so on.

True or false?

Logical values in R are written fully as TRUE and FALSE, but they are frequently abbreviated as T or F. The abbreviated version has no effect on the execution of the code, so, for example, using decreasing=T is equivalent to decreasing=TRUE in the sort function.

Logical Values

```
R> foo <- TRUE
```

```
R> foo
```

```
[1] TRUE
```

```
R> bar <- F
```

```
R> bar
```

```
[1] FALSE
```

```
R> baz <- c(T,F,F,F,T,F,T,T,T,F,T,F)
```

```
R> baz
```

```
[1] TRUE FALSE FALSE FALSE TRUE FALSE TRUE  
TRUE TRUE FALSE TRUE FALSE
```

Logical Values

A Logical Outcome: Relational Operators

Logicals are commonly used to check relationships between values. For example, you might want to know whether some number *a* is greater than a predefined threshold *b*. For this, you use the standard relational operators shown in Table 4-1, which produce logical values as results.

Table 4-1: Relational Operators

Operator	Interpretation
<code>==</code>	Equal to
<code>!=</code>	Not equal to
<code>></code>	Greater than
<code><</code>	Less than
<code>>=</code>	Greater than or equal to
<code><=</code>	Less than or equal to

Logical Values

Typically, these operators are used on numeric values:

```
R> 1==2
```

```
[1] FALSE
```

```
R> 1>2
```

```
[1] FALSE
```

```
R> (2-1)<=2
```

```
[1] TRUE
```

```
R> 1!=(2+3)
```

```
[1] TRUE
```

Logical Values

Typically, these operators are used on numeric values:

```
R> foo <- c(3,2,1,4,1,2,1,-1,0,3)
```

```
R> bar <- c(4,1,2,1,1,0,0,3,0,4)
```

```
R> length(x=foo)==length(x=bar)
```

```
[1] TRUE
```

Now consider the following four evaluations:

```
R> foo==bar
```

```
[1] FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE TRUE  
FALSE
```

```
R> foo<bar
```

```
[1] TRUE FALSE TRUE FALSE FALSE FALSE FALSE TRUE FALSE TRUE
```

```
R> foo<=bar
```

```
[1] TRUE FALSE TRUE FALSE TRUE FALSE FALSE TRUE TRUE TRUE
```

```
R> foo<=(bar+10)
```

```
[1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

Logical Values

Multiple Comparisons: Logical Operators

Logicals are especially useful when you want to examine whether multiple conditions are satisfied. Often you'll want to perform certain operations only if a number of different conditions have been met.

Table 4-2: Logical Operators Comparing Two Logical Values

Operator	Interpretation	Results
&	AND (element-wise)	TRUE & TRUE is TRUE TRUE & FALSE is FALSE FALSE & TRUE is FALSE FALSE & FALSE is FALSE
&&	AND (single comparison)	Same as & above
	OR (element-wise)	TRUE TRUE is TRUE TRUE FALSE is TRUE FALSE TRUE is TRUE FALSE FALSE is FALSE
	OR (single comparison)	Same as above
!	NOT	!TRUE is FALSE !FALSE is TRUE

Logical Values

```
R> FALSE||(T&&TRUE)||FALSE)
```

```
[1] TRUE
```

```
R> !TRUE&&TRUE
```

```
[1] FALSE
```

```
R> (T&&(TRUE||F))&&FALSE
```

```
[1] FALSE
```

```
R> (6<4)|(3!=1)
```

```
[1] TRUE
```

Logical Values

Logicals Are Numbers!

Because of the binary nature of logical values, they're often represented with TRUE as 1 and FALSE as 0. In fact, in R, if you perform elementary numeric operations on logical values, TRUE is treated like 1, and FALSE is treated like 0.

```
R> TRUE+TRUE
```

```
[1] 2
```

```
R> FALSE-TRUE
```

```
[1]-1
```

```
R> T+T+F+T+F+F+T
```

```
[1] 4
```

```
R> 1&&1
```

```
[1] TRUE
```

```
R> 1||0
```

```
[1] TRUE
```

```
R> 0&&1
```

```
[1] FALSE
```

Logical Values

Logical Subsetting and Extraction

Logicals can also be used to extract and subset elements in vectors and other objects, in the same way as you've done so far with index vectors. Rather than entering explicit indexes in the square brackets, you can supply logical flag vectors, where an element is extracted if the corresponding entry in the flag vector is TRUE.

```
R> myvec <- c(5,-2.3,4,4,4,6,8,10,40221,-8)
```

If you wanted to extract the two negative elements, you could either enter myvec[c(2,10)], or you could do the following using logical flags:

```
R> myvec[c(F,T,F,F,F,F,F,F,T)]
```

```
[1]-2.3-8.0
```

```
R> myvec<0
```

```
[1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE  
TRUE
```

```
R> myvec[myvec<0]
```

```
[1]-2.3-8.0
```

Characters

Creating a String

Character strings are indicated by double quotation marks, ". To create a string, just enter text between a pair of quotes.

```
R> foo <- "This is a character string!"
```

```
R> foo
```

```
[1] "This is a character string!"
```

```
R> length(x=foo)
```

```
[1] 1
```

```
#the total number of distinct strings
```

```
R> nchar(x=foo)
```

```
[1] 27
```

```
#Almost any combination of characters, including numbers, can be a valid character string.
```

```
R> bar <- "23.3"
```

```
R> bar
```

```
[1] "23.3"
```

Characters

Concatenation

There are two main functions used to concatenate (or glue together) one or more strings: cat and paste. The difference between the two lies in how their contents are returned. The first function, cat, sends its output directly to the console screen and doesn't formally return anything. The paste function concatenates its contents and then returns the final character string as a usable R object.

```
R> qux <- c("awesome","R","is")
```

```
R> length(x=qux)
```

```
[1] 3
```

```
R> qux
```

```
[1] "awesome" "R" "is"
```

```
R> cat(qux[2],qux[3],"totally",qux[1],"!")
```

```
R is totally awesome !
```

```
R> paste(qux[2],qux[3],"totally",qux[1],"!")
```

```
[1] "R is totally awesome!"
```

Characters

scape Sequences

```
R> cat("here is a string\nsplit\tto new\n\b\n\n\tlines")
```

here is a string
split to new
 lines

Table 4-3: Common Escape Sequences for Use in Character Strings

Escape sequence	Result
\n	Starts a newline
\t	Horizontal tab
\b	Invokes a backspace
\\\	Used as a single backslash
\"	Includes a double quote

Characters

Substrings and Matching

Pattern matching lets you inspect a given string to identify smaller strings within it.

The function `substr` takes a string `x` and extracts the part of the string between two character positions (inclusive), indicated with numbers passed as `start` and `stop` arguments.

```
R> foo <- "This is a character string!"
```

```
R> substr(x=foo,start=21,stop=27)
```

```
[1] "string!"
```